

An Empirically Calibrated Evaluation of Flood, Source-Path, Next-Hop and Gradient-Corridor Forwarding on Real Meshtastic Networks

Alen Zubic · [RelayMesh](#)

July 4, 2026 · Data window: 2026-06-27 to 2026-07-04 (UTC)

ABSTRACT. Meshtastic carries broadcast traffic — and, before firmware 2.6, all traffic — by managed flooding; 2.6 added cached next hops for direct messages, with flood as the fallback; MeshCore source-routes. This paper proposes a fourth strategy — *gradient-corridor anycast* (GradCor): one byte of state per active destination, receiver-side anycast forwarding, and **no hop limit** — termination and loop freedom follow from strict gradient descent, so a packet travels as far as real progress exists and stops only at delivery or a certified local minimum. Replayed on the **real topology of the Norwegian Meshtastic mesh** (499 nodes, 1,423 directed links reconstructed from 18,919 production traceroutes, SNR-calibrated link quality, real duty-cycle presence), the final clock-free specification **GradCor-R delivers 50.7%** of unicast packets at **21 transmissions per delivered packet** vs flooding's 25.3% at 71 tx — double the delivery at $3.3\times$ less airtime — while a MeshCore-inspired source-routing model (deliberately not a MeshCore emulation; §7) collapses to 4.0% under link asymmetry and next-hop caching reaches 30.4%. The ranking reproduces on five further production networks wherever traceroute evidence suffices. We also identify and fix a production extraction defect (traceroute responses parsed as requests) that had inflated the topology graph $2.5\times$ and understated link reciprocity (36% $\rightarrow$ 57%).

1 Introduction

LoRa mesh networks operate under one binding constraint: airtime on a shared, duty-cycled, ~kilobit-per-second channel is the only currency that matters. Meshtastic's managed flood is robust and stateless but spends the whole neighbourhood's airtime on every packet it carries; deterministic schemes promise cheaper unicast but import state that rots as nodes sleep, move and reboot. A preliminary design exploration of ours argued from first principles that a protocol should move *continuously* along the knowledge spectrum between flood (zero knowledge) and source routing (total assumed knowledge) and proposed gradient-corridor anycast: per-destination hop-distance gradients planted by a single flood, anycast forwarding down the gradient and an eligibility corridor that widens as the gradient ages — degrading gracefully toward flood instead of failing. Its supporting evidence was a synthetic simulation: 120 uniformly scattered nodes, invented link qualities and synthetic mobility churn.

Synthetic topologies flatter every protocol differently, so the natural objection is: *does the result survive a real network?* This paper answers that question using production data from [RelayMesh](#)'s ingestion of the Norwegian national Meshtastic mesh, then cross-validates on five further production networks (§6.1) — to our knowledge the first evaluation of these four strategies on fully empirical LoRa mesh topologies at this scale. Along the way the real data pushes back on the design itself: an ablation (§3.6) shows the original wall-clock staleness thresholds are unnecessary, and the paper's final specification is the simpler, clock-free form that survives it.

2 The four strategies

- **FLOOD** — Meshtastic's Managed Flood Routing^[7], the transport for all broadcast traffic and the fallback (pre-2.6: the only path) for unicast. “Managed” means the flood is disciplined, not blind: each node relays a packet at most once (duplicate suppression by packet ID) and — the mechanism

the name refers to — it first waits in an SNR-weighted contention window (worse SNR → shorter wait, so distant nodes relay first) and *cancels its own rebroadcast if it overhears another node relay the packet meanwhile*. Roles gate participation: ROUTER/REPEATER relay with priority even after hearing a duplicate, ROUTER_LATE defers to a late window and CLIENT_MUTE never relays. A hop limit (7 here, the protocol ceiling) bounds reach. Our replay models the dedup, roles and hop limit but not the overhear-cancellation — see §8, item 4.

- **PATH** — MeshCore-inspired source routing: flood-based route discovery, destination returns the recorded path over the reverse links, data then follows the strict source route with 3 retries per hop; on failure the path is invalidated and rediscovered once. PATH is a *stylized model of the source-routing mechanism, not an emulation of MeshCore*: it inherits this mesh's flood rules for discovery and runs on an organic topology with no dedicated elevated-repeater backbone — the setting MeshCore is actually designed around (§7).
- **NEXTHOP** — Meshtastic 2.6-style reactive next-hop caching for DMs: each node remembers, per destination, the relay that last delivered a packet; unicast chains down these caches and any miss falls back to a flood, which re-teaches the caches. Zero control traffic.
- **GRADCOR** — gradient-corridor anycast, the routing method this paper proposes, described in full in §3: per-destination hop-distance gradients, anycast forwarding down the gradient and an eligibility corridor that widens as knowledge decays. §3 presents both the original age-widened form (v1) and the simpler clock-free final specification (GradCor-R) that the ablation in §3.6 justifies; the evaluation runs both head-to-head.


3 The routing method: gradient-corridor anycast

This section specifies GradCor as a protocol, independently of how we evaluate it. The design answers one question: *what should a router do when its knowledge of the network is partially wrong — and it cannot know which part?* Classical answers pick an endpoint: flooding assumes nothing and pays full price every time; source routing assumes everything and breaks when any assumption fails. GradCor's answer is to make the *degree of trust in its own state* an explicit, continuously-acting protocol parameter. Every mechanism below is an instance of that one idea and Figure 1 compresses the whole method into one picture: where the design sits (a), what any single node actually does (b) and why removing the hop limit is safe (c).


We present the method as it evolved rather than only its end state, because the simplification is itself a finding. §3.1–3.5 specify the design as originally drafted (“v1”), in which trust decays with a *wall-clock age* per gradient and two fixed thresholds govern corridor widening and replanting. §3.6 then ablates those clocks on the real Norwegian network and finds them redundant: the per-hop escalation already measures staleness directly. The result is the **final specification, GradCor-R** (“reactive”) — identical machinery with every age test deleted and one byte of state per destination — which is what we recommend implementing. Readers should take v1 as the pedagogical and historical form and GradCor-R as the protocol.

GradCor at a glance: route with whatever knowledge exists, locally, for as long as progress is real

(a) The approach: uncertainty is a dial, not a failure mode



(b) One node's whole world: one number, one local rule



(c) Why hops are unlimited: progress itself is the permit




Figure 1. GradCor in one picture (shown in its final, clock-free form). **(a)** Flood and source routing are the two ends of a knowledge spectrum; GradCor is not a point on it but a slider, moved per-hop by direct failure evidence. **(b)** A node's complete view: one byte of state per active destination and five local rules evaluated on every decoded frame — eligibility is decided at the receivers, so one broadcast simultaneously probes every potential relay. **(c)** Forwarding requires the gradient to strictly decrease at every hop, so a packet may legitimately travel 12, 20, or 50 hops — a finite integer cannot descend forever, making termination and loop freedom arithmetic facts rather than counter side-effects; the flood's hop limit, by contrast, amputates at a fixed radius regardless of progress. A packet stops only at delivery or at a certified local minimum (deterministic drop + route error).

3.1 State: two numbers per destination (one, in the final spec)

Each node keeps, per *active* destination only, a single pair: the estimated hop distance to that destination (the *gradient*, g) and, in $v1$, the *age* of that estimate. Nothing else — no next-hop pointer, no path, no per-neighbour table. This is deliberate: a next-hop pointer names one specific relay and fails when that relay sleeps; a gradient only claims “I am about N hops away,” a far weaker statement that many different neighbours can satisfy. Weak claims rot slower than strong ones.

Field	Type	Meaning	Fate in final spec
g	uint8	estimated hops from this node to the destination	kept — the entire routing state
age	uint8 (hours)	time since g was planted or last refreshed	removed by the §3.6 ablation; optional replant hint only

Table 1. Per-destination state, ≈ 2 bytes in v1 and a single byte in GradCor-R, plus the node-ID key. Thirty active conversations cost a few hundred bytes — comfortably inside the RAM budget of an nRF52 target, where full link-state routing is categorically out of reach.

On the wire, a data packet carries the destination, the normal packet ID (reused for duplicate suppression), the holder's own gradient g_h and a 2-bit corridor level. That is 1–2 bytes of overhead — against a source route's full node list, this is the difference between constant and linear per-packet cost in path length.

3.2 Planting: one flood buys a whole gradient field

Gradients are created by a single flood from the destination (or equivalently, by the reverse view of any flood the destination originates — in Meshtastic practice, a NodeInfo or position broadcast can double as a planting flood at zero extra cost). Every node the flood reaches records the hop count at which it first heard it: that is the gradient, age zero. One broadcast per node, once, amortised over every subsequent packet toward that destination. There is no periodic beaconing; the protocol never spends airtime maintaining state it is not using. Figure 2 traces the two numbers through their whole life: stamped by the planting wave (a), sitting as a two-byte table row (b) and re-stamped for free whenever a delivery proves what the true working distance is (c).




Figure 2. How each node builds its two numbers. **(a)** The planting flood expands outward from the destination; every node records the hop count at which it *first* decoded it — late copies are ignored by packet-ID dedup, so g is the shortest observed distance. Nothing is requested and nothing is acknowledged. **(b)** The result at one node: a table row per active destination; the age column and its lifecycle notes are v1 — the final spec keeps only g (§3.6). **(c)** When a packet is delivered, every node on the path it actually took is re-stamped with its distance along that *proven* path and age 0 — riding the existing delivery confirmation, so active conversations keep their own gradients sharp at zero airtime cost.

One honest subtlety: the planting flood travels *outward from* the destination, while data flows *toward* it — and §4 shows 43% of real links have no observed reverse. The gradient is therefore a distance estimate measured partly over links the data cannot use in reverse: g can *underestimate* the true forward distance and a strictly-descending chain from $g=4$ can never be longer than five nodes. GradCor survives this for two reasons. First, forwarding is anycast — it never needs the specific reverse of any planting hop, only *some* decoded receiver with a lower g . Second, the corridor widens *per hop*, not only with age: a holder that finds no strictly-closer receiver retries, then admits equal- g receivers, then any receiver (Algorithm 1), so an optimistic gradient degrades into guided exploration rather than failure — and the moment one packet gets through, the delivery refresh (§3.5) re-stamps every node on the working path with its *true forward-path* distance, converting the exploration cost into a one-time fee per gradient lifetime. In the Norway replay this is measurable: 22.8% of gradient-guided deliveries took more hops than the sender's g promised (median one extra, maximum +19) and when g strictly underestimated the true

reachable distance, packets still got through at a degraded rate (16% vs 64% baseline) at roughly $6\times$ the airtime — degraded, not dead and self-correcting on first success. The evaluation in §5–6 replays this asymmetry faithfully rather than assuming it away.

3.3 Forwarding: broadcast once, let receivers race

The holder of a packet does not pick a next hop. It broadcasts once and eligibility is decided *at the receivers*: every node that decodes the frame compares its own gradient with the g_h in the header. Eligible receivers start a contention timer proportional to their gradient (ties broken by link quality); the timer that fires first forwards the packet and the other candidates hear that forward and suppress themselves. One clarification to carry through the rest of the paper: this contention timer is a *one-shot, milliseconds-scale MAC backoff armed by a packet reception* — a race resolver, not a clock. It is unrelated to the hours-scale state-aging clocks that §3.6 later removes and it is intrinsic to anycast: some tie-breaker must decide which eligible receiver goes first. This is the ExOR trick^[2], and it is what converts LoRa's broadcast medium from a liability into the protocol's main asset: one transmission simultaneously probes every potential relay, so the packet advances if *any* of them is awake and heard it — where unicast schemes stall if *the* named one didn't.

```

send(src, dst):
  if no gradient for dst at src or age  $\geq$  AGE_DEAD (24 h):
    plant(dst)                                # one flood from dst; abort if src unreached
    holder  $\leftarrow$  src; visited  $\leftarrow$  {src}

  loop:
    for level in 0..4:                        # at most 5 broadcasts per hop
      corridor  $\leftarrow$  STRICT if level  $\leq$  2 and age < AGE_WIDEN (3 h)
      EQUAL if age  $\geq$  AGE_WIDEN or level = 3
      ANY if level = 4
      broadcast(pkt, g_holder, corridor)      # one airtime unit
      # at each receiver v  $\notin$  visited that decoded the frame:
      if v = dst:                             deliver; refresh gradient along path; return
      eligible(v)  $\leftrightarrow$  g(v) < g_holder      (STRICT)
      | g(v)  $\leq$  g_holder                      (EQUAL)
      | v decoded the frame                    (ANY)
      if any eligible: winner  $\leftarrow$  contention race (min g, then best SNR); break
    if no winner: drop deterministically # local minimum; RERR toward src (§3.5)
    visited  $\leftarrow$  visited  $\cup$  {winner}; holder  $\leftarrow$  winner

```

Algorithm 1. GradCor v1 forwarding, evaluated head-to-head against its clock-free successor in §6 (constants from Table 2). In firmware the `visited` set is not carried in the packet: Meshtastic's existing packet-ID duplicate suppression plays the same role, since a node that already forwarded a packet ignores it thereafter.

3.4 The corridor: uncertainty widens eligibility, never blocks it

The corridor rule is the novel piece and the reason the protocol has no cliff-edge failure mode (Figure 3). While the gradient is fresh (<3 h), only strictly-lower-gradient receivers may forward: the packet monotonically descends toward the destination at near-unicast cost and the strict decrease on a finite integer proves both loop freedom and termination *without any hop-limit counter*. As the gradient ages past 3 h — knowledge now suspect — equal-gradient receivers are admitted, buying lateral moves around a sleeping relay or a newly-dead link at modest extra cost. If even that fails, or the gradient is very old, any receiver qualifies: the packet is now effectively flooding, which is exactly the right behaviour when the node knows nothing — and it is reached *gradually*, packet by packet, not by a mode switch that must itself be signalled. Uncertainty in GradCor is not an error state; it is a dial that trades airtime for reach in proportion to how much the state deserves to be trusted. Note that v1 drives this dial

with two triggers — wall-clock age and per-hop failure escalation — which are redundant with each other; §3.6 shows the per-hop trigger alone carries the whole effect on the real network and the final specification keeps only it.




Figure 3. The three corridor widths on a toy graph (labels show gradient g = estimated hops to destination; the source holds the packet and broadcasts once, reaching the shaded neighbourhood). The widths are identical in both forms of the protocol; only the trigger differs — per-hop failure escalation in the final spec (R), gradient age in v1. **(a)** Strict: only strictly-descending receivers (green) may forward — near-unicast cost. **(b)** Equal- g receivers admitted, allowing lateral escapes. **(c)** Any receiver eligible — including the $g=4$ neighbour that points away from the destination, which no narrower corridor admits — graceful degradation into flood.

3.5 Lifecycle: free refresh, honest death, deterministic drop

Delivery is itself the maintenance protocol. When a packet reaches the destination, every node on the path it took is stamped with its true distance-to-destination along that working path, age zero (Figure 2c; in firmware this rides the existing delivery ACK). Active conversations therefore keep their own gradients fresh *for free* and the corridor stays tight exactly where traffic flows — the state/freshness trade-off (§1) is paid only for destinations nobody is talking to. A gradient that reaches 24 h without a refresh is declared dead and the next send replants it; Figure 4 summarises the states.




Figure 4. Gradient lifecycle at one node for one destination, before and after the ablation. Top: v1 as drafted — a wall clock walks each entry through FRESH/AGED/DEAD and the corridor width is read off the state. Bottom: GradCor-R — the lifecycle collapses to two states, because corridor width is decided per hop from failure evidence instead of predicted from age. In both, the green transition is the load-bearing one: successful traffic continuously

re-stamps exactly the state it depends on, at zero airtime cost, so control traffic is only ever spent on destinations that are both wanted and quiet.

In the final specification the lifecycle therefore simplifies as the lower band of Figure 4 shows: the delivery refresh survives unchanged, the FRESH/AGED/DEAD clock states disappear because eligibility no longer consults age, and replanting degenerates to “plant if you have no gradient at all” (an optional coarse timer may evict rows that have not been refreshed for days, purely as memory hygiene).

The drop rule completes the lifecycle. A packet whose holder finds no eligible receiver at any corridor level is at a certified local minimum: the node knows nobody it can hear is closer. It drops the packet *there, deterministically* and can report a route error back toward the source (the RERR is specified but not exercised in §6's evaluation). Contrast both alternatives: a hop-limit expiry drops at an arbitrary distance-driven point regardless of progress and a stale unicast route black-holes silently. GradCor's failure is localised, observable and actionable — the property the original design brief demanded when it argued for removing the hop-limit counter, with the packet's lifetime bounded instead by the visited-set/dedup rule and the finite gradient descent.

3.6 Are the clocks necessary? An ablation

Two constants in the specification look like magic numbers: widen at 3 hours, replant at 24. A fair objection is that these are statically chosen and the soft-state literature suggests the “right” staleness timer should track the measured hazard rate of state invalidation — the reasoning behind RPL's Trickle timer^[6], which resets on observed *inconsistency* rather than on a schedule. We therefore ablated the clocks on the Norway replay (`experiment_widening.py`; Figure 5).

Three findings. **First, staleness is real:** delivery through a 1–2-hour-old gradient runs at 71%, decaying smoothly to 22% at 13–23 hours with airtime cost rising in step (Figure 5a) — a genuine hazard process with no cliff at any particular hour, so no specific threshold is “correct.” **Second, the wall-clock threshold barely matters:** sweeping the widening age from 0 (always widened) to never-widen moves delivery only between 48.0% and 52.2% (Figure 5b). **Third, the clock is redundant:** a variant that removes wall-clock age from the eligibility rule entirely — per-hop escalation only, strict → equal → any — delivers 50.9% vs the 49.7% baseline over 10 seeds, at equal cost; even deleting the replant timer too leaves it at 50.9%. The explanation is that a failed anycast broadcast is a direct, local measurement of the exact staleness the clock tries to predict: the age threshold decides in advance what the next three transmissions will discover anyway. Following the same Trickle logic, reacting to the inconsistency beats scheduling it — and the “two numbers” of §3.1 reduce to **one** (age survives only as an optional, insensitive replant hint). We also tested one adaptive variant in the opposite direction — a learned per-destination starting level that skips strict attempts a previous packet did not need — and it is *worse* (44.2%): skipping the cheap strict probes forfeits their frequent wins. The simple deterministic rule survives its own ablation; the statistical refinement does not.

This yields the final specification — what the design evolved *to*. It is Algorithm 1 with every age test deleted, nothing more:


```

# GradCor-R (final specification): Algorithm 1 with the clocks removed.
# State per active destination: g only (one byte + key). No aging tick.
send(src, dst):
    if src has no gradient for dst:
        plant(dst)                                # one flood; the only replant rule
    loop over holders as in Algorithm 1:
        for level in 0..4:
            corridor ← STRICT (g < g_holder) if level ≤ 2
                        EQUAL (g ≤ g_holder) if level = 3
                        ANY      if level = 4
            broadcast; receivers race exactly as before;
            delivery refreshes g along the path
        if no winner: drop deterministically + RERR    # unchanged

```

Algorithm 2. GradCor-R, the recommended form. Everything the evaluation praises about v1 — anycast diversity, monotone descent, no hop limit, graceful degradation, free refresh, deterministic drop — is preserved; only the prediction of staleness is replaced by its direct per-hop measurement. On the Norway replay it performs equal or better than every fixed-clock configuration — both in the isolated ablation (50.9% vs 49.7%, 10 seeds) and in the five-protocol head-to-head of Table 3 (50.7% vs 48.6%, equal cost).

Ablation: staleness is real, but a failed broadcast already measures it — the per-hop escalation makes the wall-clock thresholds unnecessary




Figure 5. Clock ablation on the Norway replay. **(a)** Outcome by gradient age at the source (baseline protocol, 3 seeds): usefulness decays smoothly — staleness is real but has no characteristic hour. **(b)** Delivery is nearly flat across the entire widening-threshold sweep (blue) and the clock-free reactive variant (green, 10 seeds) matches or beats every fixed setting: the per-hop escalation already measures staleness at the moment it matters. One caveat: this network's churn is duty-cycle-dominated with a stable router backbone; in a high-mobility mesh, wall-clock aging could recover value.

3.7 Relation to prior art and what a firmware port needs

Each ingredient is individually proven: monotone-metric descent is Babel's feasibility condition^[3] and RPL's rank rule^[4]; receiver-side anycast with contention timers is ExOR^[2]; destination-rooted gradients go back to directed diffusion in sensor networks. The composition is what is new: coupling *eligibility width to the protocol's own evidence of staleness* — wall-clock age in v1, per-hop failure escalation in the final spec — so the protocol occupies the whole flood↔unicast spectrum continuously instead of switching between brittle modes, and refreshing its state opportunistically from delivered traffic instead of periodic control packets.

A Meshtastic port is less invasive than it may sound. The firmware's router hierarchy (`src/mesh/Flood-ingRouter` → `NextHopRouter`) already contains the two hard primitives: packet-ID duplicate suppression (the visited set) and an SNR-weighted rebroadcast contention window (`RadioInterface::getTxDelayMsecWeighted`) — the managed flood already makes far nodes rebroadcast first; GradCor re-keys that same timer on (gradient, SNR). The gradient table is a small addition to per-node state alongside the NodeDB, planting piggybacks on floods that already happen (NodeInfo, po-

sition) and the corridor level fits in spare header bits. Porting GradCor-R rather than v1 removes every *scheduled* timing element: no per-entry aging tick, no periodic timer, no state clock. The one timer that remains is the contention backoff of §3.3 — one-shot, milliseconds scale, armed only by a reception — and the firmware already ships its exact skeleton: `getTxDelayMsecWeighted` maps the received SNR to a contention-window size ($-20\dots10\text{ dB} \rightarrow CW_{\min}\dots CW_{\max}$), gives ROUTER-role nodes an early window and picks a random slot inside it. GradCor-R re-keys the slot choice on (gradient, SNR) instead of SNR alone; no new timing machinery is introduced. The genuinely open engineering risk is contention-race correctness when several eligible receivers sit inside mutual capture range — duplicate forwards cost air-time but are absorbed by dedup; missed suppression is the case to instrument first (§8, item 5).

4 Data: the Norwegian mesh, one production week

All inputs were extracted from RelayMesh production systems for network “norway” (a public community network) over 2026-06-27 → 07-04:

Source	Content	Volume
Memgraph (topology graph)	:Node rows of the 7-day build window (roles, GPS; its prefix :HEARD edges were <i>not</i> used — see below)	499 nodes, 3,613 edges
ClickHouse <code>messages</code>	raw TRACEROUTE_APP packets (fwd path, <code>route_back</code> , per-link SNR)	18,919 packets
ClickHouse <code>node_activity_5min</code>	per-node hourly message counts (presence)	326 active nodes

Extracting links from traceroutes requires one correctness step the production topology builder missed at the time of this study: **31% of stored TRACEROUTE_APP rows are responses, not requests**. A response is a new packet whose envelope source and destination are swapped while `route[]` still lists the relayers in requester→responder order; parsing it as a request creates reversed edges and parsing mid-flight rows as complete paths invents endpoint hops that never happened. The corrected extraction — developed alongside this paper and since shipped to RelayMesh production — classifies each row by payload shape and hop accounting, orients both legs correctly and appends endpoints only on arrival evidence. The effect is large: the naive parse yields 4,249 directed links, the corrected one **1,423** — and measured link reciprocity rises from 36% to **57%**, i.e. much of the apparent asymmetry in the production graph is an extraction artifact, though 43% of real links still have no observed reverse. Return legs (`route_back`), previously discarded in production, contribute 303 links (~21%) seen in no other way. Role distribution: 141 CLIENT, 66 CLIENT_MUTE, 55 CLIENT_BASE, 24 ROUTER, 6 ROUTER_LATE, 1 TRACKER, 206 unset. The largest weakly connected component covers 200 nodes, the largest strongly connected component 125 and the graph is hub-dominated: the busiest routers (degree 53–107) anchor the geographic clusters.




Figure 6. The real 7-day topology after corrected extraction, plotted at GPS positions (241 nodes with valid fixes shown; 566 of 1,423 directed links have both endpoints positioned). Red/orange: ROUTER / ROUTER_LATE backbone; blue: CLIENT_BASE. Three clusters (Oslo region, Innlandet, Trøndelag) are joined by a small number of long bridge links — the same structure visible in the RelayMesh Optimization view.

5 Evaluation method

The study is a *trace-driven replay*: rather than inventing a radio environment, we reconstruct the network the production pipeline actually observed — which links existed, how good they were and when each node was awake — and then run all five forwarding protocols (the four strategies of §2, with GradCor in both its v1 and final form) over that reconstruction, hour by hour, for the same week the data describes. Figure 7 shows the pipeline end to end; the subsections below walk through each stage.




Figure 7. The trace-driven replay pipeline. Production observations (top) are reduced to three models — who can hear whom and how well, who is awake when and who is allowed to relay — which are combined into a calibrated stochastic world that the five protocols are replayed against. Note the hour quantities here (online-hours presence, ≤ 3 h gap fill, the 168-hour window) are properties of the *evaluation world* — when nodes were really awake and how long the replay runs — not protocol state; GradCor-R itself keeps no clock.


5.1 Reconstructing the link graph from traceroutes

Meshtastic's traceroute is the only mechanism in the production data that reveals *individual RF hops*. A traceroute request floods toward its destination; every node that relays it appends itself to the packet's `route[]` array and records the SNR at which it heard the previous hop in `snr_towards[]`. The responder then emits a *new packet* — envelope source and destination swapped, `route[]` copied unchanged, its own return relayers accumulating in `route_back[]` / `snr_back[]`. The true forward path of the probe is therefore `[requester] + route + [responder]` and every consecutive pair in that sequence is one real, directed, SNR-annotated radio reception (Figure 8). Three correctness rules matter (each was validated against the raw data): responses are detected by payload shape and hop accounting and their legs oriented requester→responder regardless of envelope direction; endpoints are appended only when the SNR array proves the leg actually arrived (mid-flight copies heard by a gateway must not invent the final hop); and duplicates, broadcast and placeholder IDs are dropped. The response heuristic is needed because the canonical discriminator — the protobuf `request_id` — is absent from encrypted rows ($\sim 98\%$ of traffic); on the recent unencrypted rows where it is present, the heuristic matches the canonical label on 98.7% of 871 ground-truth responses.

one TRACEROUTE_APP packet

request: `route = [R1, R2], snr_towards = [s1, s2, s3]`

reply: `route_back = [R3], snr_back = [s4, s5]`



extracted directed links: `from→R1 (s1), R1→R2 (s2), R2→to (s3) + to→R3 (s4), R3→from (s5)` ← reply links a naive forward-only parse drops

Figure 8. How one traceroute exchange becomes directed links. Solid blue: the forward leg `[requester]+route+[responder]`, whose consecutive pairs are real SNR-annotated receptions. Dashed red: the return leg from `route_back`, which travels different links when the mesh is asymmetric. Parsing both legs of all 18,919 rows with correct response orientation yields 1,423 distinct directed links; the return legs alone contribute 303 ($\sim 21\%$) that the pre-fix production builder discarded.

Aggregating across all packets, each directed link ($u \rightarrow v$) accumulates an observation count and a mean SNR. Direction is radio direction: u transmitted, v decoded. This graph is deliberately *not* symmetrized — asymmetry is a finding, not noise (§4).

5.2 Link delivery model

The trace tells us a link existed; a simulation additionally needs the probability that a single transmission on it succeeds. We model this with two multiplicative factors (Figure 9):

$$p(u \rightarrow v) = w \cdot \left[p_{\min} + (p_{\max} - p_{\min}) / (1 + e^{-(\text{SNR} - x_0)/4}) \right], \quad w = \text{obs} / (\text{obs} + k)$$

The bracket is a logistic packet-success curve in the link's mean measured SNR (raw `snr_towards` divided by 4 to approximate dB): strong links approach $p_{\max}$, weak ones fall toward $p_{\min}$. The evidence

weight w encodes a different fact: a link observed once in seven days — a ducting event, a moment of antenna alignment — is not a dependable link, however good its single SNR sample looked, while one observed hundreds of times is infrastructure. Without this term the simulator treats every fleeting link as permanently available and over-delivers dramatically at short hop counts.




Figure 9. The fitted link model. Top: per-transmission delivery probability vs link SNR, one curve per observation count — evidence, not just signal strength, determines dependability. Bottom: the distribution of mean SNR over the 1,423 real links (median -3 dB), showing most of the network's mass sits below the curve's midpoint, i.e. marginal links dominate.

Because reverse directions could be under-observed, the model includes an optional prior: an unobserved reverse link ($v \rightarrow u$) may be granted $\text{rev_p} \cdot p(u \rightarrow v)$ rather than zero, with rev_p calibrated rather than assumed. As §5.3 shows, the calibration drove rev_p to zero — once response rows are oriented correctly, the observed round-trip rate is already explained without inventing reverse links.

5.3 Calibration against two empirical anchors

The model has five free parameters (x_0 , $p_{\max}$, $p_{\min}$, k , rev_p). None are hand-picked: all were grid-searched to jointly minimise the distance to two statistics the trace itself provides.

1. **Hop-count distribution.** For 400 real (src,dst) traceroute pairs, we simulate floods at the hour each real probe was sent and compare the hop-count histogram of delivered packets with the real one (Figure 10), restricted to the 5,917 rows whose SNR array proves the forward leg actually arrived — mid-flight rows only prove a lower bound and would bias the target short. This anchors the *effective* per-link success rate along paths the network actually uses.
2. **Round-trip completion rate.** 23% of real traceroutes carry a `route_back`, i.e. completed the round trip. Simulated forward-then-reverse floods between the same pairs must approximate this rate. This anchors reverse-link generosity — exactly the quantity PATH's viability depends on.

Best fit: $x_0 = 8$ dB, $p_{\max} = 0.8$, $p_{\min} = 0.25$, $k = 10$, $\text{rev_p} = 0$. Both anchors now fit well (Figure 10; round trip 0.31 simulated vs 0.23 real). Notably the grid chose $\text{rev_p} = 0$ — with correctly-oriented links, the data is best explained by unobserved reverse directions simply *not existing*, so asymmetry is fully binding in the replay. The small remaining round-trip optimism can only *overstate* PATH (§8).




Figure 10. Calibration anchor 1. Blue: hop-count distribution of the 5,917 arrival-confirmed traceroute forward legs. Orange: hop counts of simulated floods between the same node pairs under the fitted link model. The simulator cannot reach 8 hops (its hop limit 7 binds) but reproduces the 1–7-hop mass closely.

5.4 Presence: real churn instead of synthetic mobility

The synthetic study modelled churn as random node movement. The production data says the real phenomenon is *duty cycling*: the median Norwegian node was online 22 of 168 hours and the 90th percentile 167. We therefore give every node a set of online hours — the union of its hourly activity buckets and its appearances inside any traceroute path, with gaps of ≤ 3 hours filled (a node relaying at 09:00 and 11:00 was almost certainly on at 10:00). A node that is offline in a given hour cannot receive, relay, or originate. This is a stronger and more honest stressor than mobility: state rots not because topology changed shape but because the node holding the route simply went to sleep.

5.5 Protocol implementations

All five protocols run over the identical world model; none is given information the others lack. Table 2 lists the fixed parameters. FLOOD and the discovery/fallback floods inside the other four protocols share one implementation, including a first-order MAC model: when several copies of a packet arrive at one receiver in the same flood round they collide and the strongest frame survives with capture probability 0.8 (otherwise all are lost). CLIENT_MUTE nodes never rebroadcast in any protocol.

Parameter	Value	Applies to
Hop limit	7	all floods
Per-hop unicast retries	3	PATH, NEXTHOP
Extra strict-corridor broadcast tries	2	GRADCOR v1 & R
Corridor widens at gradient age	3 h	GRADCOR v1 only
Gradient replanted at age	24 h	GRADCOR v1 only
Collision capture probability	0.8	all floods
Safety cap, transmissions per packet	2,000	all protocols

Table 2. Fixed simulation parameters (from `sim_norway.py`). The hop limit is set to Meshtastic's protocol ceiling of 7 rather than the shipped default of 3 — a choice that favours FLOOD, PATH and NEXTHOP, whose reach it bounds.

PATH discovers a route with one flood, returns the recorded path hop-by-hop over the reverse links (the discovery fails if the reply dies), then source-routes data strictly along the path with 3 tries per hop; a broken hop invalidates the route and permits one rediscovery per packet. **NEXTHOP** keeps, at every node, the relay that last delivered a packet toward each destination — learned free of charge from the parent chain of any successful flood — and chains down these caches, falling back to a flood (which re-reaches caches) on any miss or per-hop failure. **GRADCOR v1** implements Algorithm 1 exactly as specified in §3, with gradient ages advancing once per simulated hour (the upper band of Figure 4); **GRADCOR-R** implements Algorithm 2, the clock-free final specification. Both run in parallel in the head-to-head so the evolution claim of §3.6 is tested in the same arena as the competing protocols, not only in isolation. For both, the contention race is resolved ideally (lowest gradient always wins, suppression never fails) — a simplification whose consequences are discussed in §8.

5.6 Workload and metrics

Traffic replays the real demand pattern: 30 (source, destination) pairs sampled from the actual traceroute pair-frequency distribution (weighted, without replacement), restricted to endpoints online ≥ 20 hours in the week so conversations can exist at all. Each pair sends one packet per hour, but only in hours where both endpoints were genuinely online — a message to a sleeping node is not a routing failure and must not be scored as one. Ten seeds give $\approx 30,000$ packet attempts per protocol. Every broadcast or unicast attempt costs one transmission; gradient-planting, discovery and fallback floods are charged to the packet that triggered them, so no protocol externalises its control traffic. We report delivery rate, transmissions per packet sent (offered-load cost) and transmissions per packet *delivered* (the airtime price of a success — the currency that matters, per §1).

6 Results

Protocol	Delivery %	Tx / packet sent	Tx / packet delivered
FLOOD (hop limit 7)	25.3	17.9	71.0
PATH (stylized source routing, §2)	4.0	18.9	473.2
NEXTHOP (2.6-style cache)	30.4	18.2	59.9
GRADCOR v1 (age-widened corridor)	48.6	10.4	21.4
GRADCOR-R (final spec, no clocks)	50.7	10.8	21.3

Table 3. Aggregate results over 10 seeds, 30 real pairs, 168 real hours, with both GradCor forms evaluated in parallel. Absolute percentages are model outputs; the ranking and the ratios are the robust findings. Per-seed delivery varies with the sampled traffic pairs (s.d. 4.5–5.4 points), but the *paired* per-seed differences are tight: GradCor-R beats FLOOD on all ten seeds (+19.5 to +29.9 points, mean +25.3, s.d. 3.4), while R – v1 is $+2.1 \pm 3.0$ — a genuine tie within noise, slightly ahead, independently confirming the §3.6 ablation in the full head-to-head.




Figure 11. Delivery rate (bars, left axis) and airtime cost per delivered packet (diamonds, right axis, log scale), with both GradCor forms side by side. GradCor-R dominates flooding on both axes simultaneously: +25.4 points of delivery (2.0 $\times$) at 3.3 $\times$ less airtime per delivered packet — with no state clocks.




Figure 12. Delivery rate across the real week (6-hour buckets); light green is v1, dark green GradCor-R — the two track each other throughout (R converges slightly faster in the cold-start hours, before v1's aged gradients earn their first refresh), and the advantage over flooding is stable through daily presence cycles. PATH never recovers because its routes break faster than they amortize.

6.1 Does it generalise? Five more production networks

A single network, however real, is one draw from the space of mesh topologies. To test whether the result is a property of the protocol or a property of Norway, we ran the identical harness on the five other production networks with the most traceroute traffic in the same week: Bay Area Mesh (3,637 nodes), Florida Mesh, SoCal Mesh, Meshtastic Portugal and Meshtastic Italia. Everything was recomputed from scratch per network — link graph, presence, roles and crucially the *link-model calibration*, re-fitted against each network's own hop-count distribution and round-trip completion rate (the fitted parameters genuinely differ: xo ranges 0–12 dB and the reverse-link prior stays at zero everywhere except Florida and SoCal). Reusing Norway's constants elsewhere would have been methodologically wrong; nothing about the protocols themselves was changed or tuned.




Figure 13. Cross-network validation on six production networks (5 seeds each, same 168-hour week, per-network calibration). Top: delivery rate; bottom: airtime per delivered packet (log). On every network with adequate traceroute evidence the ranking is identical — GradCor-R $\approx$ GradCor v1 > NEXTHOP > FLOOD >> PATH — with GradCor-R delivering 1.8–2.5 $\times$ flood's rate at 2.5–3.7 $\times$ less airtime per delivered packet. Italia (starred) is the honest boundary case: its 2,492 traceroutes yield only 146 usable links for 2,207 nodes, gradients have almost nothing to descend on and gradient routing degrades to flood level while NEXTHOP's simpler cache wins.

Network	Nodes	Links	Recip %	FLOOD %	GRADCOR-R %	R vs flood	R tx/dlv
Norway	499	1,423	57	25.5	51.2	2.0 $\times$ @ 3.3 $\times$ cheaper	21.3
Meshtastic PT	1,208	4,147	36	24.0	43.2	1.8 $\times$ @ 3.5 $\times$ cheaper	25.1
SoCal Mesh	1,351	3,461	31	13.2	27.8	2.1 $\times$ @ 3.7 $\times$ cheaper	35.2
Bay Area Mesh	3,637	6,895	32	5.6	12.9	2.3 $\times$ @ 3.4 $\times$ cheaper	54.3
Florida Mesh	1,394	3,994	37	4.7	11.9	2.5 $\times$ @ 2.5 $\times$ cheaper	46.1
Meshtastic Italia	2,207	146	30	6.2	5.1	0.8 $\times$ — below eligibility bar	39.9

Table 4. Per-network summary (full five-protocol data in Figure 13 and `multinet_results.csv`). Norway's row is its 5-seed re-run within this harness, hence the small differences from Table 3's 10-seed values — a useful incidental read on seed noise ($\approx \pm 0.5$ points). Absolute delivery scales with evidence density — the dense US networks have far more nodes per usable link, so every protocol delivers less there — but the *ratios* hold. Italia's row is the operative precondition, not a counterexample: gradient routing needs enough traceroute-derived link evidence to plant gradients on — exactly the evidence the extraction fixes of §7 (now shipped) increase.

Two observations deserve emphasis. First, the ranking and the ratios — the two things §8 argues are the trustworthy outputs of this methodology — survive a 7 $\times$ range in network size, a 30–57% range in link

reciprocity and independently fitted link models; the Norway result is not a topology accident. Second, the one failure is diagnostic rather than embarrassing: where traceroute evidence is too thin to define gradients (Italia, 0.07 links per node), GradCor degrades to roughly flood performance instead of below it — the graceful-degradation property doing exactly what it was designed to do at the extreme — while the evidence-free NEXTHOP cache takes the lead. Evidence density is thus a measurable eligibility criterion for deploying gradient routing on a given mesh.

7 Discussion

The real topology strengthens the synthetic result. On the synthetic mesh GradCor matched flood's delivery at $2.5\text{--}3\times$ less airtime; on the real Norwegian mesh — sparser and harsher than the naive extraction suggested — GradCor-R *doubles* flood's delivery while keeping the airtime ratio. The mechanism is visible in the graph structure: the ROUTER backbone (degree 53–107 hubs) is present nearly 24/7, so gradients through it stay valid for many hours, while anycast forwarding exploits the receiver diversity that broadcast gives for free — any of several hub-adjacent nodes can carry the packet downhill, so no single lossy link is load-bearing. Notably, the corrected (sparser) graph *widened* the margin: when links are scarce and marginal, spending them precisely matters more. And the head-to-head confirms the evolution claim in vivo: the clock-free final spec runs even with — here slightly ahead of — the age-widened original (50.7% vs 48.6% at equal cost), with two fewer constants and half the state.

Source routing fails for a measurable, structural reason. PATH needs every hop of a recorded path to work in the named direction and needs the reverse path to deliver the route reply. With 43% of links lacking any working reverse and median 3–4 hop paths, the probability that a full source route remains valid is small and each failure triggers a discovery flood *plus* a reply that often dies on an asymmetric link. Its 473 tx per delivered packet is worse than never routing at all. This is the algorithm stripped of MeshCore's real-world advantage — deployments built around fixed elevated repeaters with engineered symmetric links — so it should be read as “source routing on an organic mesh,” not as a verdict on MeshCore in its intended setting.

Next-hop caching under-delivers under duty-cycle churn. NEXTHOP edges out flooding (30.4% vs 25.3% at slightly lower cost) but captures little of GradCor's gain: with the median node online 13% of the time, per-destination caches go stale between conversations and most packets fall back to the teaching flood anyway. Its one virtue — zero control traffic and trivial state — remains, which is presumably why Meshtastic 2.6 shipped it.

Production findings. Building the replay surfaced two extraction defects in the production topology builder, one of them substantial. First, it parsed traceroute *responses* as requests, misorienting their edges and inventing endpoint hops from mid-flight rows: the corrected extraction shrinks the Norway graph from 3,613 to 1,423 directed links and raises measured reciprocity from 36% to 57% — meaning the old graph both overstated connectivity and overstated asymmetry, distorting any coverage, bridge, or role-recommendation rule computed on it. Second, it discarded `route_back/` `snr_back`, which contribute $\sim 21\%$ of real links and the only direct per-link bidirectionality evidence — the single most decision-relevant link property for any routing improvement. Both defects were fixed and shipped to production during this work; because topology windows are ephemeral and rebuilt from raw rows, every subsequent build self-heals with no data migration.

8 Threats to validity

1. **Residual round-trip optimism.** Simulated round-trip completion (0.31) still slightly exceeds the observed rate (0.23), so reverse-path behaviour is modelled a little generously even with `rev_p = 0`. This can only inflate PATH; a tighter model would push its 4.0% lower. FLOOD, NEXTHOP and GRADCOR never require a named reverse hop and are insensitive to it.

2. **Survivorship bias.** Traceroutes reveal only links that worked at least once in the window; usable-but-unexercised links are missing, which slightly disadvantages every protocol equally and understates true connectivity.
3. **Observation bias.** Ground truth comes from MQTT gateways; RF events no gateway heard are invisible. Absolute delivery rates are therefore model outputs, not measurements — only rankings and ratios should be quoted.
4. **MAC simplification.** Same-round collision with 0.8 capture is crude; there is no inter-packet airtime contention or duty-cycle budget. Real broadcast storms are collision-dominated, so this likely still flatters FLOOD, i.e. the real gap would widen in GradCor's favour — but that remains unproven here, as it was in the synthetic pilot. One simplification cuts the other way: the replay's baseline flood omits Managed Flood Routing's overhear-cancellation (a node cancelling its queued rebroadcast when it hears another node relay first), which understates flood's real airtime efficiency — while GradCor's analogous suppression is granted ideally (item 5). We measured this bias by re-running the head-to-head with a cancellation-aware flood (contention-ordered: routers first, clients worst-SNR-first, ROUTER_LATE last; `experiment_cancellation.py`): cancellation saves flood 15% of transmissions but also *costs* it 2.0 points of delivery ($25.3 \rightarrow 23.3\%$) — cancelled rebroadcasts are sometimes the only path to edge nodes, the coverage hole ROUTER_LATE exists to patch. Net effect on the headline: GradCor-R's delivery advantage widens slightly ($2.0\times \rightarrow 2.1\times$) and its airtime-per-delivered advantage compresses modestly ($3.3\times \rightarrow 3.0\times$). No conclusion changes. The headline tables deliberately keep the baseline flood: the link model was calibrated so baseline floods reproduce real traceroute statistics, and those real floods already ran *with* cancellation — its aggregate effect is partially absorbed into the fitted link probabilities, so layering explicit cancellation on top without recalibrating would double-count the suppression. This run is therefore a bias bound, not a better model.
5. **Contention-timer realism.** GradCor's anycast suppression (“lowest gradient wins, others overhear and mute”) is assumed perfect. Near-simultaneous forwarders within capture range are the mechanism most likely to misbehave on real radios. Note this is the one timing mechanism the §3.6 ablation could *not* remove — it is a reception-armed race resolver, not a state clock — which makes it the entire remaining timing surface of GradCor-R and the first thing to prototype in firmware.
6. **Free, lossless gradient refresh.** The replay re-stamps every node on a delivered packet's path instantly and at zero cost; in firmware the refresh rides the delivery confirmation, which travels real links and can be lost — particularly on asymmetric paths. This optimism is GradCor-specific. Its consequence is bounded — a missed refresh only leaves gradients staler, costing exploration rather than creating a black hole, and the $v1 \leftrightarrow R$ equivalence suggests low sensitivity to exactly how staleness is handled — but the refresh delivery rate belongs on the firmware-prototype checklist alongside contention.
7. **Independent link trials and no end-to-end retries.** Every transmission is an independent Bernoulli draw from the fitted link probability; real LoRa links fail in correlated bursts (fading, temporary obstructions), which penalises rapid per-hop retries more than receiver diversity, and no protocol was granted end-to-end retransmissions (a real Meshtastic DM retries several times). Both simplifications apply to all five protocols equally, but they mean absolute delivery rates should not be compared with field ACK statistics.
8. **Unicast scope.** Every claim in this paper concerns unicast traffic (DMs, ACKs, traceroute-class exchanges). Broadcast traffic — positions, telemetry, channel messages, the majority of load on many meshes — is inherently flood-shaped and is neither helped nor harmed by GradCor; the replay also carries no background broadcast load contending for airtime. The airtime savings shown here apply to the unicast fraction of a mesh's traffic only.

9 Conclusions and next steps

On a real 499-node national mesh with empirically calibrated links and real duty-cycle churn, gradient-corridor anycast in its final clock-free form delivered twice the *unicast* packets of managed flooding at 40% less per-packet airtime and $3.3\times$ less airtime per delivered packet, while source routing was structurally non-viable and next-hop caching was only marginally better than flooding. The ordering matches the synthetic study's prediction; the corrected, sparser real topology *widened* the margin. And the ordering is not Norway's: five further production networks, independently calibrated, reproduce it wherever traceroute evidence is dense enough to plant gradients — with the one evidence-starved network degrading gracefully to flood level rather than failing and defining a measurable eligibility criterion (usable links per node) for deployment. The same replay also simplified the design itself: the clock ablation (§3.6) showed the wall-clock thresholds the method started with are redundant and the head-to-head (Table 3) confirms it — the specification the paper ends with, GradCor-R (Algorithm 2), matches the original with one byte of state per active destination and one deterministic escalation ladder: no state clocks, no tuning constants beyond retry counts and the millisecond contention window it inherits unchanged from the existing managed flood and no statistical estimation anywhere in the forwarding path. What began as “two numbers and two thresholds” ends as “one number and one rule.”

Recommended next steps, in order of information value:

1. Close the decoder-engine `request_id` gap (extract it at ingest, optionally redecode historical rows): this upgrades the response classification of §5.1 from a 98.7%-accurate heuristic to canonical labels. The extraction fix itself (response orientation, arrival-gated endpoints, `route_back` ingestion) already shipped to production during this work.
2. Cross-validate reverse-link behaviour with NEIGHBORINFO edges where available and re-fit against the 23% round-trip anchor.
3. Add concurrent-traffic airtime contention to test the collision-dominated regime.
4. Prototype the anycast contention race on hardware (two RAK4631s plus one Heltec V3 within capture range) before any firmware investment in the full protocol. With the state clocks gone, this millisecond backoff is GradCor-R's *entire* timing surface and the primitive to extend already exists (`RadioInterface::getTxDelayMsecWeighted`, re-keyed on gradient); the experiment to run is suppression reliability: how often does the losing candidate fail to hear the winner and forward a duplicate?

References

1. Meshtastic firmware 2.6 release notes — next-hop routing for direct messages.
2. S. Biswas, R. Morris, “ExOR: Opportunistic Multi-Hop Routing for Wireless Networks,” SIGCOMM 2005 (anycast contention forwarding).
3. J. Chroboczek, “The Babel Routing Protocol,” RFC 8966 (feasibility condition / monotonic metric descent).
4. T. Winter et al., “RPL: IPv6 Routing Protocol for Low-Power and Lossy Networks,” RFC 6550 (rank rule).
5. [RelayMesh](#) production data, Norway network, 2026-06-27–07-04. Reproduction package (data exports and all simulation, experiment and figure code: `sim_norway.py`, `experiment_asymmetry.py`, `experiment_widening.py`, `experiment_multinet.py`, `experiment_cancellation.py`, `fig_*.py`; cross-network exports under `data_networks/`) is published alongside this paper.
6. P. Levis, T. Clausen, J. Hui, O. Gnawali, J. Ko, “The Trickle Algorithm,” RFC 6206 (suppress on consistency, react to inconsistency — not to a schedule).
7. Meshtastic documentation: “[Mesh Broadcast Algorithm](#)” and “[Why Meshtastic Uses Managed Flood Routing](#).”